

Rohan Desai

Boston, MA | rohan.desai@example.com | (617) 555-0127 | linkedin.com/in/rohandesai-example | github.com/rohandesai-example | rohandesai.dev

EDUCATION

Northeastern University

Boston, MA | Master of Science in Computer Software Engineering | Expected May 2026 | GPA: 4.0/4.0

Relevant coursework: Object-Oriented Design, Algorithms, Web Design, Cloud Computing, Design Patterns

University of Mumbai

Mumbai, India | Bachelor of Technology in Information Technology | May 2023 | GPA: 3.6/4.0

Relevant coursework: Data Structures, System Design, Operating Systems, DevOps, Distributed Systems, Database Management Systems

SKILLS

- **Programming Languages:** Python, Go, Java, JavaScript, HTML, SQL, Protobuf
- **Web Technologies:** React, Next.js, Express, Node.js, Tailwind CSS, TypeScript, GraphQL
- **Databases:** Oracle, MongoDB, MySQL, PostgreSQL, Redis, NoSQL
- **Cloud and DevOps:** DigitalOcean, AWS, GCP, Azure, Docker, Kubernetes, Jenkins, Terraform
- **Libraries and Frameworks:** NumPy, pandas, TensorFlow
- **Tools:** Git, Linux, Jira, Confluence, Bitbucket, Jasmine, Grafana, Storybook, Jest

EXPERIENCE

DigitalOcean | Software Engineer Intern

Boston, MA | Jun 2025 - Sep 2025

- Improved DigitalOcean Marketplace, which offers preconfigured applications and services deployable on cloud servers.
- Built a production-grade LLM inference service powering Marketplace Search in Go with gRPC, Protobuf, Redis, and RAG, reducing p90 query latency from 3.2 seconds to 1.5 seconds over eight weeks.
- Designed evaluation pipelines with a golden dataset, increasing click-through rate by 15% and deployments from search by 10% during the internship period.
- Improved internal RPCs and enhanced the Vendor Portal and Pre-deploy service to streamline Marketplace partner onboarding and integration.
- Created and deployed custom Marketplace Droplets with shell scripting and Packer, automating provisioning and reducing deployment setup time from hours to minutes.
- Implemented monitoring and automated testing pipelines for Marketplace Droplets, reducing release-related failures and improving deployment reliability.

HDFC Bank | Software Engineer

Mumbai, India | Jul 2023 - Jul 2024

- Scaled a digital payments and mobile wallet platform by optimizing REST APIs with Swagger/OpenAPI, Node.js, and Java, supporting 35% higher transaction throughput after optimization.
- Architected a real-time chat system using microservices and WebSockets, rolled out to thousands of concurrent users.
- Containerized services with Docker, deployed via Kubernetes, and implemented CI/CD pipelines that reduced deployment time by 25%.
- Collaborated across teams to identify performance bottlenecks, improving system responsiveness and engagement across millions of active users.

University of Mumbai | Research Assistant

Mumbai, India | Jul 2022 - Jan 2023

- Conducted EEG signal research to detect seizures by building real-time, high-frequency data pipelines in Python.
- Used wavelet transforms, CNN-LSTM models, multithreading, feature extraction, ICA, and bandpass filtering, improving ML classifier accuracy to 85%.

PROJECTS

NotePressAI

Go, TypeScript, AWS, PostgreSQL, React, LLM | Jan 2026 - Mar 2026

- Built a full-stack workspace unifying document editing, diagramming, and AI assistance with Go and React.
- Implemented six block types, eight canvas tools, Mermaid syntax import, and deployment on Railway and Netlify.
- Engineered AI flow validation detecting dead ends, loops, missing branches, and unreachable nodes with 95% accuracy.
- Optimized a context-aware assistant, reducing API costs by 80% while keeping response time under four seconds.

CloudScale

Go, Terraform, AWS, PostgreSQL, CI/CD, GitHub Actions | Sep 2025 - Nov 2025

- Architected auto-scaling AWS infrastructure with Terraform, including ALB, multi-AZ RDS PostgreSQL, Route 53 DNS, KMS encryption, and auto-scaling groups across dev and demo environments.
- Built a Go REST API with Gin and GORM, including authentication, S3 storage, SNS/Lambda email verification, and CI/CD pipelines with zero-downtime deployments through GitHub Actions.

Inference Router

Python, SDK, LLM Optimization | Apr 2026 - May 2026

- Built a Python SDK that routes LLM requests across model tiers using pluggable complexity, cost, and latency strategies.
- Supported AWS Bedrock, OpenAI, and OpenAI-compatible APIs through a unified provider interface, reducing inference costs by routing simple queries to smaller models.